

Clustering Analysis^{*}

Ke Chen

To carry out this assessed coursework, you will need the Python notebook, [clustering.ipynb](#), which contains code to get you started with each of the assignments, along with several data files required by different assignments (see the related assignments for details of those data sets) in this coursework. All above are available in a zipped file on BlackBoard alongside this document.

Clustering analysis in its own right is a classical topic in unsupervised learning to learn a high-level summary of data, which provides a representation of data from a global perspective. In this coursework, you are asked to implement state-of-the-art clustering algorithms in Python and apply your implementations along with the provided Python implementation of clustering algorithms to several meaningful synthetic datasets for clustering analysis.

1 K -means Clustering Analysis

K -means is one of the most commonly used clustering analysis algorithm. In this work, you are asked to apply the provided K -means functions to synthetic datasets to understand its behaviour and how to use a cluster validity index to decide the proper number of clusters underlying a data set.

Assignment 1 [3 marks] Apply the built-in K -means function, [sklearn.cluster.KMeans](#), in the [scikit-learn](#) library with Euclidean distance to the 2-D dataset of 3 clusters, [kmeans_data_1.npy](#), with three different initial conditions given in the notebook, [clustering.ipynb](#). In your answer notebook, (a) implement the function, [partition](#), that produces a partition for a given dataset so that the data points in this dataset are grouped into clusters based on their closest centroids. All the points in each cluster are assigned the same label specified by their corresponding centroid; (b) apply the built-in K -means function to the dataset mentioned above for clustering analysis with 3 given different initialisation settings, respectively; and (c) based on the clustering analysis results achieved in (b), use your [partition](#) function implemented in (a) to create 6 scatter plots corresponding to 3 initial partitions and 3 final partitions in order to visualise 6 partitions. You are asked to display 6 plots in a 2×3 grid; the first row shows 3 initial partitions and the second row shows 3 final partitions aligned with their corresponding initial partitions shown in the first row. In each scatter plot, you must mark the cluster centroids with the red-coloured "+", all the data points in a cluster must be marked with the same colour (but different from red to allow for seeing their centroid clearly) and different clusters must be indicated by different colours. (**Hint:** to carry out the display format described above, you may use the built-in function, [matplotlib.pyplot.subplot](#).)

Assignment 2 [3 marks] K -means algorithm cannot be used until the hyperparameter K (the number of clusters) is set up. In a real application, however, the number of clusters is often un-

^{*} **Assessed Coursework:** the deadline and requirements can be found at the end of this document.

known. In this circumstance, the scatter-based **F-ratio index**¹ may be applied to decide the number of clusters. In your answer notebook, (a) implement the scatter-based F-ratio index in Python where Euclidean distance is used; (b) for $K = 2, 3, \dots, 10$, run the K -means built-in function, `sklearn.cluster.KMeans`, in the `scikit-learn` library with Euclidean distance on the dataset, `kmeans_data_2.npy` (for each K , you must run K -means with **3** different random initialisation conditions set by yourself), then plot F-ratio index (y-axis) versus K (x-axis) (in this plot, for each K , use only the **least** F-ratio index value measured on 3 partitions resulting from different initialisation conditions), and report the optimal number of clusters in this data set; and (c) display the final partition corresponding to the optimal number of clusters you find out in (b) with the the same display format described in **Assignment 1**.

2 Spectral Clustering Analysis

As a state-of-the-art clustering analysis method, spectral clustering can deal with data of nonlinear and non-convex manifold with the connectivity criteria. In this work, you are asked to implement a spectral clustering algorithm in Python and apply your own implementation to synthetic datasets in order to understand how spectral clustering works on data sets in different nature.

Assignment 3 [4 marks] In your answer notebook, implement the **asymmetric normalised spectral clustering**² algorithm where the **fully connected graph** is used to generate the Laplacian matrix used in spectral clustering. You are asked to use two functions: `gaussian_similarity` for generating the fully connected graph and `asymmetric_SC` for the asymmetric normalised spectral clustering in your implementation. (**Hint:** To implement the `asymmetric_SC` function, you can use the built-in function, `np.linalg.eig`, in the `numpy` library for eigen analysis and two built-in functions in the `scikit-learn` library, `sklearn.cluster.KMeans` for K -means clustering and `pairwise_distances` for measuring the distances between data points.)

Assignment 4 [3 marks] Apply your implemented `gaussian_similarity` function in **Assignment 3** to the dataset, `SC_data_1.npy`, to generate a Laplacian matrix. In your answer notebook, (a) describe how you find out an appropriate hyperparameter value in the Gaussian kernel and explicitly report this value; (b) conduct eigen analysis on the Laplacian matrix, list all the elements of the eigenvector corresponding to the second smallest eigenvalue or the 1st smallest non-zero eigenvalue in a 4×10 table carried out by a proper 2-D array; and (c) with the result achieved in (b), provide the computational evidence to decide how many clusters in this dataset and display the resultant partition in a 2-D plot where different clusters must be marked in different colours.

Assignment 5 [2 marks] Apply your implemented `gaussian_similarity` and `asymmetric_SC` functions in **Assignment 3** to the dataset, `SC_data_1.npy`, with the number of clusters you find out in **Assignment 4**, and two other datasets of two clusters, `SC_data_2.npy` and `SC_data_3.npy`, respectively. In your answer notebook, (a) display the 3 resultant partitions of the datasets in 2-D plots where clusters must be marked in different colours, and (b) check whether this clustering algorithm works for all 3 datasets. If not, identify the dataset(s) it fails and explain why.

¹For the definition of the F-ratio index, see “ K -means Clustering” lecture note.

²For the algorithmic description of this algorithm, see “Spectral Clustering” lecture note.

3 Hierarchical Clustering Analysis

As one of the most important clustering analysis techniques, hierarchical clustering enables one to achieve all possible clusters at different levels for a data set. In this work, you are asked to apply the commonly-used agglomerative algorithm with different cluster-distance metrics to a synthetic dataset in order to understand how this algorithm works.

Assignment 6 [4 marks] You are asked to use Euclidean distance to produce the initial distance matrix and further apply the built-in function, `scipy.cluster.hierarchy` with three different cluster-distance metrics³, single-linkage, complete-linkage and group-average, in the `scipy` library to the dataset, `HC_data.npy`, respectively. In your answer notebook, (a) plot three dendrogram trees achieved by the use of three cluster-distance metrics in the agglomerative algorithm; (b) implement the function, `get_longest_lifetime`, used to find out the longest K -cluster lifetime from a dendrogram; and (c) use your `get_longest_lifetime` function implemented in (b) to report the longest K -cluster lifetimes along with their corresponding K (number of clusters) found in terms of the longest K -cluster lifetime criterion for 3 dendrograms achieved from (a) and further display 3 partitions corresponding to the longest K -cluster lifetimes. For visualisation, you must arrange 3 plots, each with the explicit cluster-distance metric title, in one row in the order of single-linkage, complete-linkage and group-average from left to right.

4 Ensemble Clustering Analysis

As one of the state-of-the-art clustering analysis techniques, ensemble clustering algorithms allow for combining different types of clustering algorithms and working on various feature sets of a data set to reach a synergy for clustering analysis. In this work, you are asked to implement the evidence-accumulated clustering algorithm and apply your implementation to two benchmark data sets to understand how this algorithm works in practice.

Assignment 7 [2 marks] In your answer notebook, implement a function, `ensemble_clustering`, for the evidence-accumulated clustering algorithm⁴ based on K -means and the agglomerative algorithms in Python where K -means algorithm with Euclidean distance is used to generate initial partitions. (**Hint:** To implement the `ensemble_clustering` function, you can use the built-in functions, `sklearn.cluster.KMeans`, in the `scikit-learn` library for K -means clustering, the `scipy.cluster.hierarchy.linkage` and `scipy.spatial.distance.squareform` in the `scipy` library for hierarchical clustering.)

Assignment 8 [4 marks] Apply your `ensemble_clustering` function implemented in **Assignment 7** to two datasets, `SC_data_2.npy` and `SC_data_3.npy`, respectively. In your experiment, you are asked to figure out how to decide how many initial partitions need to be generated by K -means algorithm with Euclidean distance and what an appropriate cluster-distance metric is to make the clustering ensemble algorithm work properly on a given dataset. In your answer notebook, (a) for **each** of two datasets, describe how you generate initial partitions and what the cluster-distance metric that lead to satisfactory results with justification (you need to state the number of initial

³For the definition of cluster-distance metrics, see “Hierarchical Clustering” lecture note.

⁴For the algorithmic description of this algorithm, see “Clustering Ensemble” lecture note.

partitions and the chosen cluster-distance metric explicitly), and (b) plot two dendrogram trees produced by the cluster ensemble algorithm on two datasets and further display the optimal consensus partitions of two dataset with the longest K -cluster lifetimes achieved, respectively, in the format of 2-D scatter plot where clusters in a partition must be marked in different colours.

Requirement: Before starting working on this assessed coursework, you need to

1. download all the files required by this coursework from Blackboard as specified at the beginning of this document;
2. unzip the file then you should be see a Jupyter notebook file named `clustering.ipynb` and one sub-directory named `Data` (you must keep this directory/sub-directory structure and its name unchanged when you work on this coursework);
3. rename `clustering.ipynb` in the directory as `yourfullname_clustering.ipynb`. For instance, if your name is “John Smith”, your filename should be `john_smith_clustering.ipynb`. This file will be your answer notebook to be submitted for marking, so you must include everything required by the coursework in this Jupyter notebook.

Deliverable: Only your answer notebook, `yourfullname_clustering.ipynb`, which should include all your code, output, answers and your interpretation/justification. In this Jupyter notebook, all assignments have been separated with the clear delimiters. You must put your stuff regarding an assignment in those cells related to this assignment and, if necessary, create new cells within the delimiters of this assignment.

Your answer notebook, `yourfullname_clustering.ipynb`, must be submitted via the Blackboard.

Marking: Marking is on the basis of (1) correctness of results and quality of comments on your code; (2) rigorous experimentation; (3) how informative and clear your description/answers presented in your answer book; and (4) your knowledge exhibited, interpretation and justification.

Late Submission Policy: The default university late submission policy (for details, see the official document: <http://documents.manchester.ac.uk/DocuInfo.aspx?DocID=29825>) is applied to this coursework.

Extension Policy: The default departmental extension policy is applied to this coursework. That is, no extension is allowed unless you have a mitigating circumstance. If you have any mitigating circumstance and want to make an extension, you should submit the completed mitigating circumstance form to SSO (for details, see the departmental Mitigating Circumstances page: <http://studentnet.cs.manchester.ac.uk/assessment/mitigatingcircumstances.php?view=pgt>). Note: the decision will be made by the departmental mitigating circumstance panel rather than the lecturer.

Deadline: 18:00 GMT, 16th December 2021 (Thursday)